

# Introduction to Computer Networks



## Reliable Transmission

© All rights reserved. No part of this publication and file may be reproduced, stored in a retrieval system, or transmitted in any form or by any means, electronic, mechanical, photocopying, recording or otherwise, without prior written permission of Professor Nen-Fu Huang (E-mail: [nfhuang@cs.nthu.edu.tw](mailto:nfhuang@cs.nthu.edu.tw)).

# Outline


---

- **Introduction**
- **Stop-and-Wait protocol**
- **Sliding Window Protocol**
- **Issues with Sliding Window Protocol**

# Reliable Transmission

---

- Data transmissions **over a communication link** may have errors caused by signal interference.
- Usually, CRC is used to detect errors.
- Some error codes are strong enough to correct errors but the overhead is typically too high.
- Corrupt frames must be discarded.
- For reliable transmission, we must recover from these discarded frames.



# Reliable Transmission

---

- We can do this by using a combination of two fundamental mechanisms
  - **Acknowledgements**
  - **Timeouts**
- An **acknowledgement** (**ACK** for short) is a small control frame that a protocol sends back to its peer saying that it has received the earlier frame.
  - A **control frame** is a frame with header only (no data).
- The receipt of an **acknowledgement** indicates to the sender that its frame was successfully delivered.

# Reliable Transmission

---

- If the sender does not receive an *acknowledgment* after a reasonable amount of time, then it retransmits the original frame.
- The action of waiting a reasonable amount of time is called a *timeout*.
- The general strategy of using *acknowledgements* and *timeouts* to implement reliable delivery is sometimes called **Automatic Repeat reQuest (ARQ)**.

# Outline

---


- Introduction
- **Stop-and-Wait protocol**
- Sliding Window Protocol
- Issues with Sliding Window Protocol

# Stop and Wait Protocol

---

- Idea of **stop-and-wait protocol** is straightforward
  - After transmitting one frame, the sender **waits for an acknowledgement** before transmitting the next frame.
  - If the acknowledgement does not arrive after a certain period of time, the sender **times out and retransmits** the original frame

# Stop and Wait Protocol




Four different scenarios for the stop-and-wait algorithm.

(a) The ACK is received before the timer expires;

(b) The original frame is lost;



# Stop and Wait Protocol



Four different scenarios for the stop-and-wait algorithm.

(c) The ACK is lost;

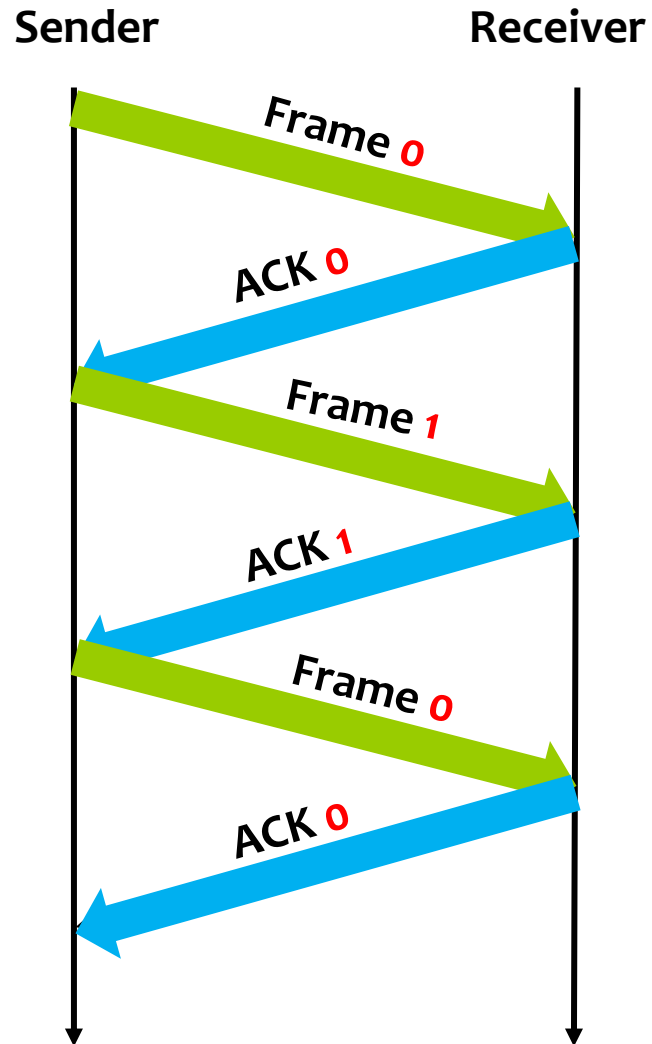
(d) The timeout fires too soon (or the ACK is delayed)

# Stop and Wait Protocol

---

- If the acknowledgment is **lost or delayed** in arriving
  - The sender times out and retransmits the original frame.
  - As a result, **duplicate copies of frames will be delivered**
- How to solve this ?
  - **Use 1 bit sequence number (0 or 1)**
  - When the sender **retransmits frame 0**, the receiver can determine that this is a second copy of frame 0 rather than the first copy of frame 1 and therefore can ignore it (the receiver still acknowledges it, in case the first acknowledgement was lost)

# Stop and Wait Protocol



Timeline for stop-and-wait with 1-bit sequence number

# Stop and Wait Protocol

---


- The sender has **only one outstanding frame** on the link at a time -- Far below the link's capacity
- For example, consider a 2 Mbps link with a 40 ms RTT
  - The link has a delay  $\times$  bandwidth product of 80 Kb or **8 KB**
  - Since the sender can send only one frame per RTT and assuming **a frame size of 1 KB**
  - Maximum Sending rate
    - ▶ Bits per frame/Time per frame =  $8\text{kb} \div 40\text{ms} = 200 \text{ Kbps}$   
Or **about 1/10 of the link's capacity (2Mbps)**
  - To fully use the link, the sender should transmit up to ten frames before having to wait for an ACK.

# Outline

---

- Introduction
- Stop-and-Wait protocol
- **Sliding Window Protocol**
- Issues with Sliding Window Protocol

# Sliding Window Protocol



Timeline for Sliding Window Protocol

# Sliding Window Protocol


---

- Sender assigns a sequence number denoted as **SeqNum** to each frame.
  - Assume it can grow infinitely large
- Sender maintains three variables
  - **Sending Window Size (SWS)**
    - ▶ Upper bound on the number of outstanding (unacknowledged) frames that the sender can transmit
  - **Last Acknowledgement Received (LAR)**
    - ▶ Sequence number of the last acknowledgement received
  - **Last Frame Sent (LFS)**
    - ▶ Sequence number of the last frame sent

# Sliding Window Protocol

- Sender also maintains the following invariant

$$LFS - LAR \leq SWS$$



Sliding Window on Sender



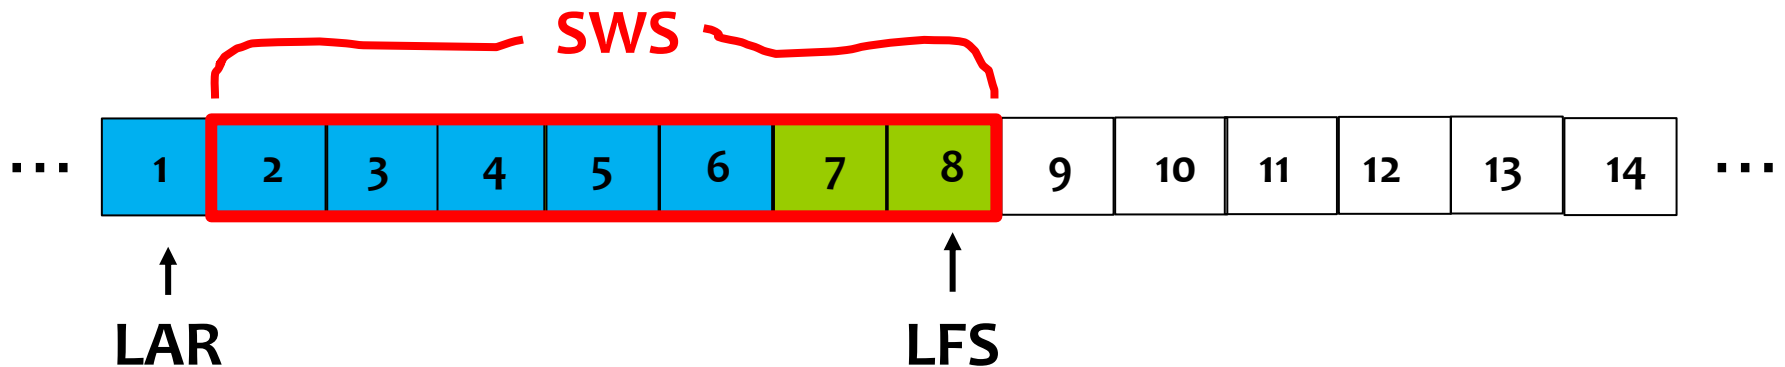
# Sliding Window Protocol

---

- When an acknowledgement arrives
  - the sender **moves LAR to right**, thereby allowing the sender to transmit another frame
- Also the sender associates a **timer** with each frame it transmits
  - It retransmits the frame if the timer expires before the ACK is received
- Note that the sender needs to **buffer** up to SWS frames for retransmissions, if necessary

# Sliding Window Protocol

- When the sender is allowed to slide its Window ?
  - When the acknowledgement of  $LAR+1$  is received



Sliding Window on Sender

# Sliding Window Protocol


---

- Receiver maintains three variables
  - **Receiving Window Size (RWS)**
    - ▶ Upper bound on the number of out-of-order frames that the receiver is willing to accept
  - **Largest Acceptable Frame (LAF)**
    - ▶ Sequence number of the largest acceptable frame
  - **Last Frame Received (LFR)**
    - ▶ Sequence number of the last frame received

# Sliding Window Protocol

- Receiver also maintains the following invariant

$$\text{LAF} - \text{LFR} \leq \text{RWS}$$



Sliding Window on Receiver

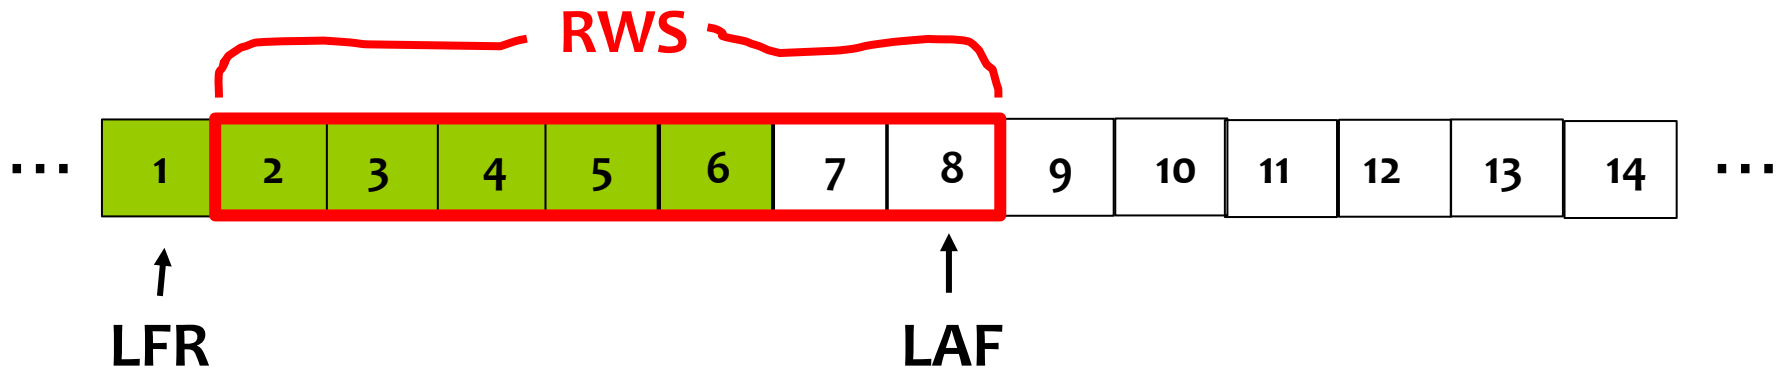
# Sliding Window Protocol

---

- When a frame with sequence number **SeqNum** arrives, what does the receiver do?
  - If  **$\text{SeqNum} \leq \text{LFR}$  or  $\text{SeqNum} > \text{LAF}$** 
    - ▶ Discard it (**the frame is outside the receiver window**)
  - If  **$\text{LFR} < \text{SeqNum} \leq \text{LAF}$** 
    - ▶ Accept it

# Sliding Window Protocol

- When the receiver is allowed to slide its Window ?
  - When the frame of  $LFR+1$  is received



Sliding Window on Receiver

# Sliding Window Protocol

---

- For example, suppose **LFR = 1** and **RWS = 7**  
    **→ LAF = 8**
- If frames 4,6,3 and 5 arrive sequentially, they are out of order, but will still be buffered because they are within the receiver window
- But no ACK will be sent since frame 2 is yet to arrive
- Finally, frame 2 arrives (retransmitted or delayed)
- Now Receiver **Acknowledges Frame 6**  
    **and bumps LFR to 6**  
    **and LAF to 13 (Window sliding)**

# Sliding Window Protocol


---

## ■ Cumulative Acknowledgement

- Let **SeqNumToAck** denote the largest sequence number not yet acknowledged, such that all frames with sequence number **less than SeqNumToAck have been received**
- The receiver acknowledges the receipt of **SeqNumToAck** even if high-numbered packets have been received
  - ▶ This acknowledgement is said to be **cumulative**.
- The receiver then sets
  - ▶  **$LFR = SeqNumToAck - 1$**  and
  - ▶  **$LAF = LFR + RWS$  (Window sliding)**



# Sliding Window Protocol



Cumulative Acknowledgement example (**SeqNumToAck = 2**)

# Outline

---

- Introduction
- Stop-and-Wait protocol
- Sliding Window Protocol
- **Issues with Sliding Window Protocol**

# Issues with Sliding Window Protocol

---

- When **timeout occurs**, the amount of data in transit decreases
  - Since the sender is unable to advance its window
- When the **packet loss occurs**, this scheme is no longer keeping the pipe full
  - The longer it takes to notice that a packet loss has occurred, the more severe the problem becomes
- How to improve this
  - **Negative Acknowledgement (NAK)**
  - **Additional Acknowledgement**
  - **Selective Acknowledgement**

# Issues with Sliding Window Protocol

---

## ■ Negative Acknowledgement (NAK)

- Receiver sends NAK for frame 2 when frame 3 arrive (in the cumulative acknowledgement example)
  - ▶ However this is unnecessary since sender's timeout mechanism will be sufficient to catch the situation

## ■ Additional Acknowledgement

- Receiver sends additional ACK for frame 2 when frame 3 arrives
  - ▶ Sender uses duplicate ACK as a clue for frame loss

## ■ Selective Acknowledgement

- Receiver will acknowledge exactly those frames it has received, rather than the highest number frames
  - ▶ Receiver will acknowledge frames 3,4,5, and 6
  - ▶ Sender knows frame 2 is lost
  - ▶ Sender can keep the pipe full (additional complexity)

# Issues with Sliding Window Protocol

---

## How to select the window size ?

- SWS is easy to calculate
  - ▶  $\text{Delay} \times \text{Bandwidth}$
- RWS is more flexible
  - ▶ Two common settings
    - »  $RWS = 1$ 

No buffer at the receiver for frames that arrive out of order
    - »  $RWS = SWS$ 

The receiver can buffer frames that the sender transmits
- Does it make any sense to keep  $RWS > SWS$  ?

# Issues with Sliding Window Protocol

---

## ■ Finite Sequence Number

- Frame sequence number is specified in the header field

- ▶ Finite size

- » 3 bits: eight possible sequence number: 0, 1, 2, 3, 4, 5, 6, 7

- ▶ It is necessary to wrap around

# Issues with Sliding Window Protocol

---

- How to distinguish between different frames of the same sequence number?
  - Number of possible sequence number must be larger than the number of outstanding frames allowed
    - ▶ Stop and Wait: One outstanding frame
      - » 2 distinct sequence number (0 and 1)
    - ▶ Let **MaxSeqNum** be the number of available sequence numbers
    - ▶  $SWS + 1 \leq \text{MaxSeqNum}$ 
      - Is this sufficient?
      - Depends on RWS
      - If  $RWS = 1$ , then sufficient
      - If  $RWS = SWS$ , then not good enough

# Issues with Sliding Window Protocol

---

For example, we have eight sequence numbers


0, 1, 2, 3, 4, 5, 6, 7

$$RWS = SWS = 7 = 8-1$$

1. Sender sends 0, 1, ..., 6
  2. Receiver receives 0, 1, ... ,6
  3. Receiver acknowledges 0, 1, ..., 6  
but **ACK (0, 1, ..., 6) are lost**
  4. Sender timeouts and retransmits **0, 1, ...,5, 6**
  5. Receiver is expecting 7, **0, ...., 5**
- ➔ frames 0-5 will be accepted (duplicated !!)  
frame 6 will be discarded (correct)



# Issues with Sliding Window Protocol



Problem when  $SWS + 1 \leq \text{MaxSeqNum}$  and  $SWS = RWS$   
 $\text{MaxSeqNum} = 8, SWS = RWS = 7$

# Issues with Sliding Window Protocol

---

To avoid this,

If  $RWS = SWS$

$$SWS < (MaxSeqNum + 1)/2$$

# Issues with Sliding Window Protocol

---

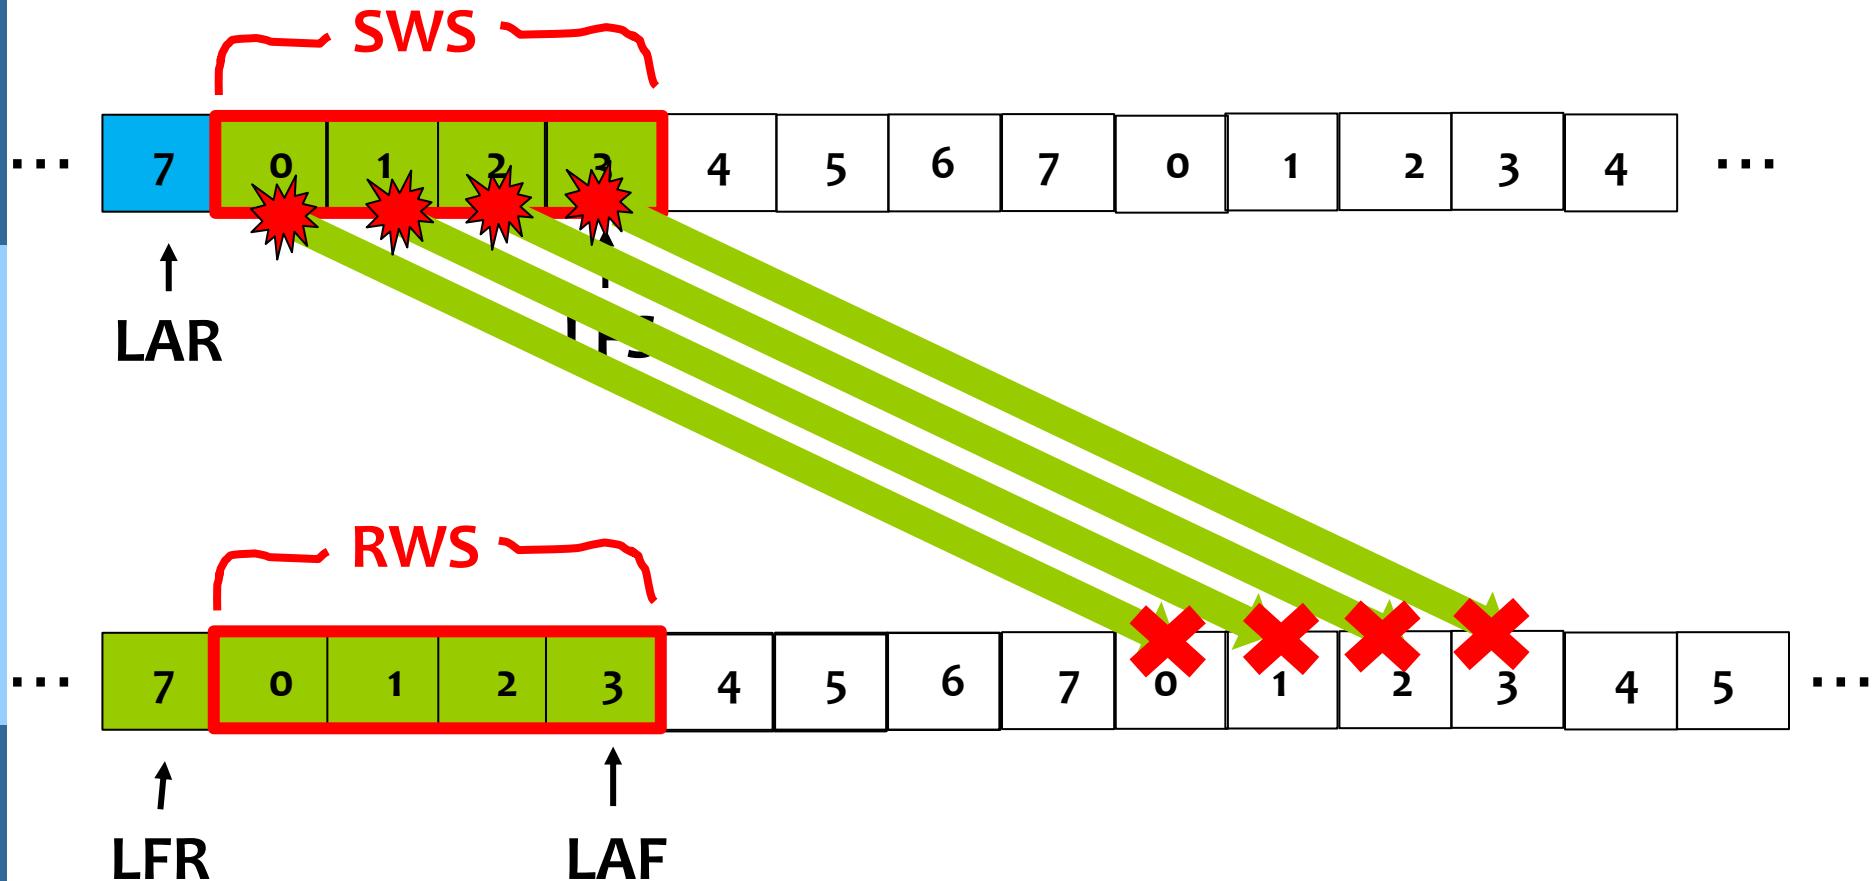
For example, we have eight sequence numbers

0, 1, 2, 3, 4, 5, 6, 7

$$RWS = SWS = 4 < (8+1)/2 = 4.5$$

1. Sender sends 0, 1, 2, 3
  2. Receiver receives 0, 1, 2, 3
  3. Receiver acknowledges 0, 1, 2, 3  
but **ACK (0, 1, 2, 3) are lost**
  4. Sender timeouts and retransmits **0, 1, 2, 3**
  5. Receiver is expecting 4,5,6,7
- ➔ **all frames 0-3 will be discarded (correct now !!)**

# Issues with Sliding Window Protocol



Example when  $SWS < (MaxSeqNum+1)/2$   
 $MaxSeqNum = 8, SWS = RWS = 4 < (8+1)/2$

# Issues with Sliding Window Protocol

---

- Sliding Window Protocol provides three features
  - **Reliable Transmission**
  - **Preserve the order**
    - ▶ Each frame has a sequence number
    - ▶ Out of order frames will be buffered
  - **Flow control**
    - ▶ Receiver is able to throttle the sender by setting the value of RWS
    - ▶ Keeps the sender from overrunning the receiver

# Summary

---

- Two fundamental mechanisms are used to provide reliable transmission over a communication link
  - Acknowledgements
  - Timeouts
- Stop-and-Wait Protocol is a reliable protocol, but the performance is not good enough
  - Only one outstanding frame
  - Receiver may receive duplicated frames

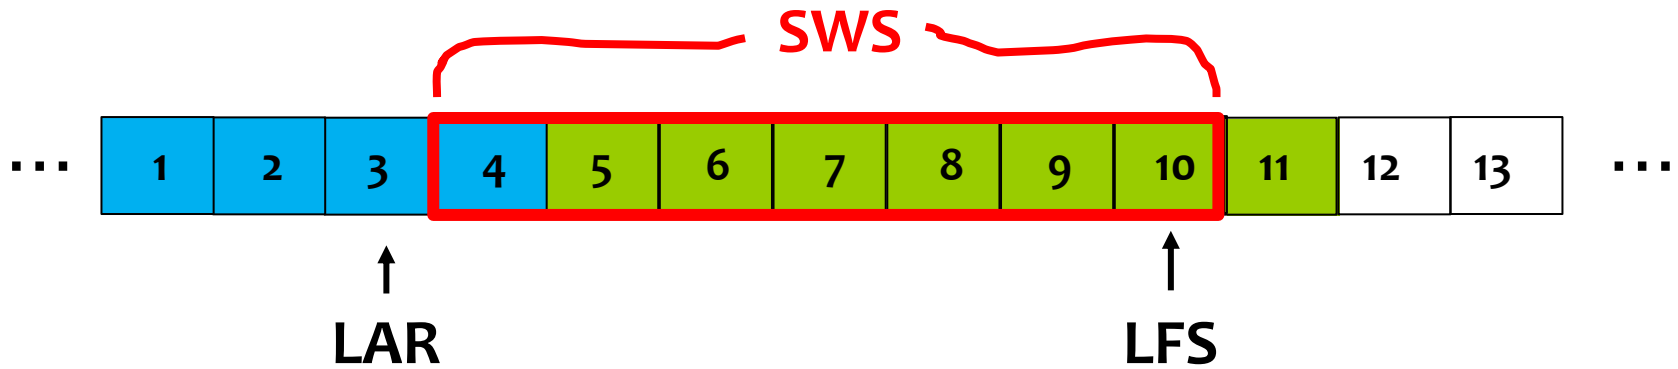
# Summary

---

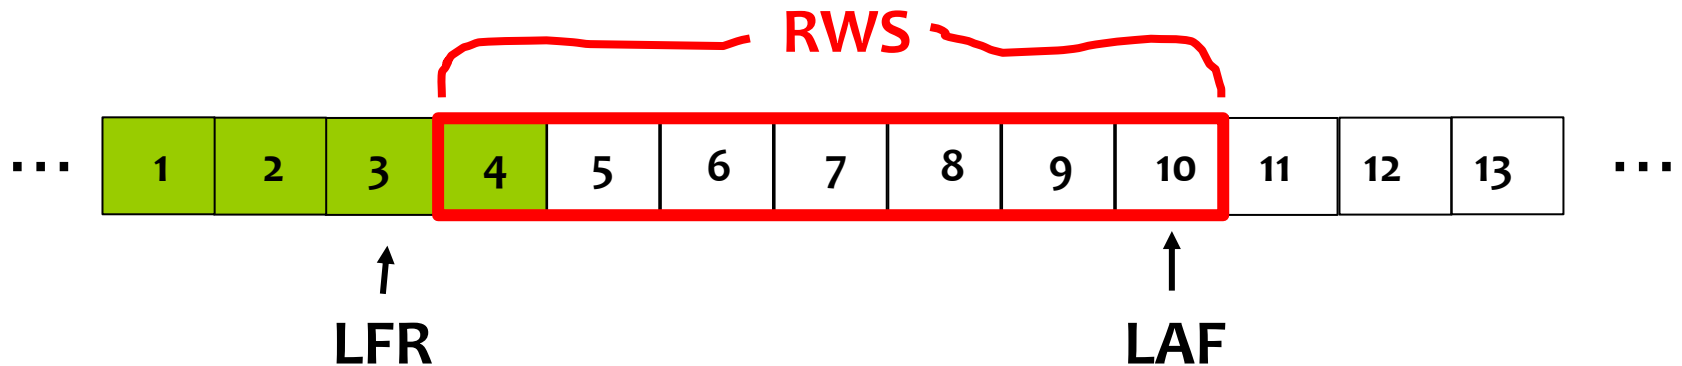
- **Sliding Window Protocol is a reliable and good performance protocol**
  - Sequence number is added for each frame
  - Multiple outstanding frames (keeping pipe full)
  - **Sender**
    - ▶ **Sending Window Size (SWS) = delay x bandwidth**
    - ▶ **Last Acknowledgement Received (LAR)**
    - ▶ **Last Frame Sent (LFS)**
  - **Receiver**
    - ▶ **Receiving Window Size (RWS) = 1 or SWS**
    - ▶ **Largest Acceptable Frame (LAF)**
    - ▶ **Last Frame Received (LFR)**

# Summary

- Sender maintains the invariant:  $LFS - LAR \leq SWS$



- Receiver maintains the invariant:  $LAF - LFR \leq RWS$





# Summary

---

- The values of SWS and RWS may be different.
- If  $RWS = SWS$ , it is better  $SWS < (MaxSeqNum + 1)/2$  otherwise, the receiver may still receive duplicated frames.
- The Acknowledgement mechanism is an option
  - NAK
  - Cumulative acknowledgement
- Sliding Window Protocol provides three features
  - Reliable Transmission
  - Preserve the order
  - Flow control (receiver determines the RWS)